

Technical Architecture & Governance of TML Smart Contracts

1. The Enforcement Primitive (Software Reality)

The foundational layer of the Ternary Moral Logic (TML) framework is the Enforcement Primitive, a smart contract architecture designed to function as a deterministic State-Transition Engine. This engine operationalizes ethical commitments by embedding them directly into the execution logic of a smart contract, which acts as a gatekeeper for all subsequent actions on an EVM-compatible platform. The core principle is that every transaction or state-altering function call must first pass through a TML evaluation. This evaluation does not produce a simple boolean (true/false) but a tri-state output (+1, 0, -1), which dictates the subsequent state transition of the governed contract. This design moves beyond traditional access control lists or role-based permissions, introducing a more nuanced, logic-driven governance model where the "constitution" is encoded into the contract itself. The architecture is built to be strict, efficient, and resistant to administrative override, ensuring that the ethical framework is not merely advisory but is the primary determinant of execution.

1.1. Smart Contract as a State-Transition Engine

The TML smart contract is fundamentally a transaction-based state machine, a concept central to the Ethereum Virtual Machine (EVM) itself. Ethereum's global state, encompassing all account balances and contract storage, transitions to a new state with each processed block of transactions. The TML contract leverages and extends this model by creating a localized, application-specific state machine whose transitions are governed by the output of a TML evaluation function. This function, $Y(S, T) \rightarrow S'$, takes the current contract state S and a transaction T (or a function call) as input and produces a new state S' based on the TML verdict. The determinism of this function is paramount; given the same initial state and transaction input, every node in the network must compute the exact same resulting state to maintain consensus. This ensures that the ethical enforcement is not only predictable but also verifiable by all participants in the network. The contract's logic is therefore designed to be a pure function in this regard, avoiding any non-deterministic elements like reliance on block timestamps or external, unverified data sources.

1.1.1. Core TML Logic Implementation

The core of the TML enforcement primitive is the evaluation logic that processes an incoming transaction and produces a tri-state output. This logic is designed to be a formal, verifiable system, potentially modeled using a logic framework such as Past-Time Linear Temporal Logic (PLTL), which has been proposed for expressing and validating complex transaction trace

properties in real-time . While the specific axioms of TML are defined by the overarching framework, their implementation within the smart contract must be a series of deterministic checks. For instance, a transaction could be evaluated against a set of rules that check for compliance with principles derived from human rights documents or environmental protocols. These checks would analyze the transaction's parameters, its potential effects on the contract's state, and its position within a sequence of transactions (its trace). The output of this evaluation is not a simple "allow" or "deny" but one of three distinct values, each triggering a different state transition within the contract's lifecycle. This approach allows for a more sophisticated response than a simple binary gate, introducing a mechanism for handling ambiguity and ethical uncertainty.

The implementation of this logic must be meticulously designed to be gas-efficient, as complex computations can be prohibitively expensive on-chain. The research into instrumenting safety properties within smart contracts highlights a significant challenge: such runtime guards can be gas-intensive, increasing the cost for users and potentially rendering them impractical for on-the-fly validation . Therefore, the TML logic should be optimized, potentially by pre-computing parts of the decision tree off-chain and storing only the necessary verification data on-chain. The contract would then use cryptographic proofs, such as Merkle proofs, to verify that a transaction complies with the pre-computed rules without having to execute the entire logic on-chain. This hybrid approach balances the need for on-chain enforcement with the practical constraints of the EVM's computational model. The formalism of the logic, whether based on PLTL or another system, provides a rigorous foundation for ensuring that the contract's behavior is consistent, predictable, and aligned with the intended ethical framework.

1.1.2. Solidity Interfaces and Constants

To ensure clarity, interoperability, and formal verification, the TML enforcement primitive is defined through a set of Solidity interfaces and constants. These components serve as the public API and the immutable parameters of the ethical enforcement layer.

Solidity Interfaces:

The primary interface, `ITMLEnforcer`, defines the functions that any contract wishing to be governed by TML must implement or interact with. This standardization allows for a modular and composable ecosystem of TML-compliant applications.

```
// SPDX-License-Identifier: TML-Constitutional-Code
pragma solidity ^0.8.19;

/// @title ITMLEnforcer
```

```
/// @notice The interface for the Ternary Moral Logic enforcement primitive.
interface ITMLEnforcer {
    /// @notice Represents the tri-state output of the TML evaluation.
    /// @dev Using an enum for type safety and clarity.
    enum TMLVerdict {
        REFUSE, // -1
        REVIEW_HOLD, // 0
        PROCEED // +1
    }

    /// @notice Represents the finite states of the governed contract.
    enum ContractState {
        STATE_ACTIVE,
        STATE_REVIEW_HOLD,
        STATE_FLAGGED,
        STATE_INTEGRITY_FROZEN
    }

    /// @notice Emitted when a transaction is evaluated by the TML engine.
    /// @param transactionHash The hash of the transaction being evaluated.
    /// @param verdict The TML verdict for the transaction.
    /// @param logId A unique identifier for the evaluation log entry.
    event TMLEvaluation(bytes32 indexed transactionHash, TMLVerdict verdict, bytes32 logId);

    /// @notice Emitted when the contract's state changes.
    /// @param previousState The state before the transition.
    /// @param newState The state after the transition.
    /// @param triggerReason The reason for the state change (e.g., TML verdict, council action).
    event StateTransition(ContractState previousState, ContractState newState, string triggerReason);
```

```

    /// @notice Evaluates a transaction against the TML logic.
    /// @dev This is the core function that must be called before any
    state-changing operation.
    /// @param data The calldata of the transaction to be evaluated.
    /// @return verdict The TMLVerdict enum value.
    function evaluateTransaction(bytes calldata data) external
returns (TMLVerdict verdict);

    /// @notice Gets the current state of the contract.
    /// @return state The current ContractState enum value.
    function getContractState() external view returns (ContractState
state);

    /// @notice Gets the "Always Memory" root hash.
    /// @return rootHash The Keccak-256 hash of the Merkle root of
the off-chain logs.
    function getAlwaysMemoryRoot() external view returns (bytes32
rootHash);
}

```

Constants:

The following constants define the immutable parameters of the TML system, ensuring that the rules are fixed at deployment and cannot be altered by any administrator, thus enforcing the "No God Mode" principle.

```

// SPDX-License-Identifier: TML-Constitutional-Code
pragma solidity ^0.8.19;

/// @title TMLConstants
/// @notice A library containing the immutable constants for the TML
enforcement layer.

```

```

library TMLConstants {
    /// @dev The duration a contract can remain in STATE REVIEW HOLD
before auto-failing.
    /// 72 hours = 72 * 60 * 60 = 259,200 seconds.
    uint256 public constant REVIEW_HOLD_DURATION = 259200;

    /// @dev The Keccak-256 hash of the mandated corpora (e.g., Human
Rights documents).
    /// This is a placeholder; the actual hash would be determined at
deployment.
    bytes32 public constant CONSTITUTION_HASH =
0x1234567890abcdef1234567890abcdef1234567890abcdef1234567890abcdef;

    /// @dev The number of council members required to resolve a
`STATE REVIEW HOLD`.
    uint256 public constant COUNCIL_CONSENSUS_THRESHOLD = 5;

    /// @dev The total number of members in the Stewardship Council.
    uint256 public constant COUNCIL_TOTAL_MEMBERS = 7;

    /// @dev The maximum number of consecutive `REFUSE` verdicts
before a contract is `STATE FLAGGED`.
    uint256 public constant FLAGGING_THRESHOLD = 10;
}

```

These interfaces and constants provide a clear, auditable, and standardized foundation for the TML enforcement primitive. They define the language through which contracts interact with the ethical layer and establish the unchangeable rules that govern its behavior, forming the basis of the "Constitutional Code."

1.1.3. State Machine Design and Implementation

The TML enforcement primitive is implemented as a Finite State Machine (FSM), a well-established model in computer science for designing systems with a limited number of distinct states and defined transitions between them. This model is particularly well-suited for blockchain applications, as it ensures deterministic and predictable behavior, which is essential

for consensus. The FSM within the TML contract governs its own lifecycle, transitioning between states based on the TML evaluation of incoming transactions and the actions of the Stewardship Council. The design of this FSM is critical to the system's security and integrity, as it must handle not only the "happy path" of approved transactions but also the complex scenarios of rejection, escalation, and potential system compromise. The implementation must be robust, with clear trigger conditions for each state transition and safeguards against invalid or malicious state changes.

The FSM is defined by its states, transitions, and the logic that governs those transitions. The states are explicitly defined in the `ContractState` enum within the `ITMLEnforcer` interface: `STATE_ACTIVE`, `STATE REVIEW HOLD`, `STATE_FLAGGED`, and `STATE_INTEGRITY_FROZEN`. The transitions between these states are triggered by specific events. For example, a TML verdict of `PROCEED (+1)` while the contract is `STATE_ACTIVE` will result in the transaction being executed, and the contract remains in `STATE_ACTIVE`. Conversely, a `REFUSE (-1)` verdict will cause the transaction to revert, also leaving the contract in `STATE_ACTIVE`. The most critical transition is triggered by the `REVIEW_HOLD (0)` verdict, which moves the contract from `STATE_ACTIVE` to `STATE REVIEW HOLD`. This state is a mandatory pause, designed to handle ethical ambiguity or uncertainty. The contract can only exit this state through a definitive action, either by a consensus decision from the Stewardship Council or by the expiration of a time-bound auto-fail mechanism. This strict, rule-based approach to state transitions is what gives the TML framework its power and reliability, ensuring that ethical considerations are not bypassed.

1.2. Tri-State Enforcement Mapping

The core innovation of the TML enforcement primitive is its tri-state logic, which moves beyond the binary "allow/deny" paradigm of traditional access control. This mapping defines how the contract responds to the evaluation of a transaction, with each of the three possible outputs—`+1` (Proceed), `-1` (Refuse), and `0` (Sacred Zero)—triggering a distinct and deterministic action. This nuanced approach allows the system to handle not only clear-cut cases of compliance and violation but also the gray areas of ethical uncertainty that are often encountered in complex systems. The mapping is designed to be strict and unambiguous, ensuring that every transaction is met with a clear and predictable response. This is crucial for maintaining the integrity of the system and for providing a reliable framework for developers and users to interact with. The tri-state logic is the heart of the TML enforcement primitive, and its careful design and implementation are essential to the success of the entire framework.

The enforcement of this mapping is handled by the `evaluateTransaction` function, which acts as the entry point for all state-changing operations. This function is responsible for running the transaction through the TML evaluation logic and then executing the corresponding action based on the verdict. The implementation of this function must be robust and secure, with careful attention to gas efficiency and reentrancy protection. The function's output is not just a boolean but a member of the `TMLVerdict` enum, which is then used by the contract's main

logic to determine the next step. This design ensures that the tri-state logic is deeply embedded into the contract's execution flow, making it an integral part of the system's behavior. The mapping is not just a set of rules but a fundamental aspect of the contract's architecture, shaping how it interacts with the world and how it enforces the ethical principles it is designed to uphold.

1.2.1. +1 (Proceed): Transaction Execution

A TML evaluation resulting in a `+1` (Proceed) verdict signifies that the transaction has been deemed ethically compliant and is authorized for execution. This is the "happy path" for a TML-governed contract, where the transaction passes all the checks and balances of the TML framework without any ambiguity or conflict. When the `evaluateTransaction` function returns `TMLVerdict.PROCEED`, the calling function is free to execute its intended logic, whether it's a token transfer, a state update, or any other operation. The contract remains in its current state, typically `STATE_ACTIVE`, and the transaction is processed as normal. This outcome represents a clear and unambiguous approval, providing a high degree of confidence to both the user and the system that the action is in line with the established ethical guidelines. The `PROCEED` verdict is the most straightforward of the three outcomes, but its significance lies in the rigorous evaluation that precedes it, ensuring that only truly compliant transactions are allowed to proceed.

The implementation of the `PROCEED` logic is relatively simple from a state machine perspective. The key is to ensure that the evaluation leading to this verdict is thorough and reliable. The TML logic must be designed to cover all relevant ethical considerations, and the evaluation process must be deterministic and verifiable. The `TMLEvaluation` event is emitted with the `PROCEED` verdict, providing a transparent and auditable record of the decision. This log entry is crucial for off-chain analysis and for building trust in the system. The simplicity of the `PROCEED` path is a feature, not a bug; it ensures that the common case of compliant transactions is handled efficiently and without unnecessary overhead. The complexity of the system is focused on the other two verdicts, which handle the more challenging scenarios of rejection and uncertainty. By making the `PROCEED` path as clean and efficient as possible, the TML framework minimizes the cost and friction for legitimate use cases, encouraging adoption and fostering a healthy ecosystem.

1.2.2. -1 (Refuse): Transaction Reversion

A TML evaluation resulting in a `-1` (Refuse) verdict indicates that the transaction has been found to be in clear violation of the TML's ethical axioms. This is a definitive rejection, and the contract's response is to revert the transaction, preventing any state changes from occurring. When the `evaluateTransaction` function returns `TMLVerdict.REFUSE`, the calling function must immediately revert, typically using Solidity's `revert()` statement. This ensures that the malicious or non-compliant action is blocked, and the system remains in a safe and consistent state. The gas spent on the transaction up to the point of reversion is still consumed, which

serves as a deterrent against spam or repeated attempts to violate the rules. The `REFUSE` verdict is a critical security feature of the TML framework, providing a robust defense against actions that are fundamentally incompatible with the system's ethical principles. It is a clear and unambiguous "no," leaving no room for negotiation or override.

The implementation of the `REFUSE` logic is also straightforward from a state machine perspective. The contract remains in its current state (e.g., `STATE_ACTIVE`), and the transaction is simply blocked. However, the system can be designed to track repeated `REFUSE` verdicts. If a contract or an address accumulates a certain number of refusals within a given period, it could trigger a transition to the `STATE_FLAGGED` state. This would impose additional restrictions, such as higher gas costs or a requirement for manual review, providing an additional layer of protection against persistent bad actors. The `TMLEvaluation` event is emitted with the `REFUSE` verdict, creating a permanent and public record of the violation. This transparency is essential for accountability and for allowing the community to monitor the system's security. The `REFUSE` verdict is a powerful tool for enforcing the "hard rules" of the TML framework, ensuring that the ethical boundaries are not crossed.

1.2.3. 0 (Sacred Zero): Triggering `STATE REVIEW HOLD`

The `0` (Sacred Zero) verdict is the most nuanced and critical component of the TML tri-state logic. It is triggered when the TML evaluation cannot produce a definitive `+1` (Proceed) or `-1` (Refuse) verdict. This situation arises when the transaction presents an ethical ambiguity, a novel scenario not covered by existing axioms, or a conflict between different ethical principles. The `Sacred Zero` is not a permanent freeze or a denial; rather, it is a mandated pause, a "time-out" for the system to seek further guidance. When the `evaluateTransaction` function returns `TMLVerdict.REVIEW_HOLD`, the contract immediately transitions to the `STATE REVIEW HOLD` state. In this state, all state-changing operations are blocked, and the contract awaits a resolution. This is the primary mechanism for escalating complex ethical decisions to the human oversight body, the Stewardship Council. The `Sacred Zero` is a crucial feature that acknowledges the limitations of purely algorithmic decision-making and provides a structured pathway for human intervention.

The implementation of the `Sacred Zero` logic is more complex than the other two verdicts. It involves a state transition to `STATE REVIEW HOLD` and the initiation of a timer. The contract must remain in this state for a predefined duration (e.g., 72 hours), during which the Stewardship Council can review the case and provide a verdict. If the council reaches a consensus within the time limit, the contract transitions to a new state based on that verdict. If the time expires without a resolution, the contract may automatically transition to a fail-safe state, such as `STATE_FLAGGED` or even `STATE_INTEGRITY_FROZEN`, depending on the severity of the ambiguity. The `TMLEvaluation` event is emitted with the `REVIEW_HOLD` verdict, and a unique `logId` is generated to link the on-chain event to the off-chain review process. This ensures a clear and auditable trail of the entire escalation and resolution process. The `Sacred`

`Zero` is the bridge between the deterministic world of smart contracts and the nuanced, deliberative world of human ethics, making it the cornerstone of the TML governance model.

1.3. Sequence Diagrams

To illustrate the operational flow of the TML enforcement primitive, the following sequence diagrams depict the interactions between a user, the TML-governed smart contract, the TML evaluation engine, and the Stewardship Council. These diagrams provide a visual representation of the different paths a transaction can take, from the "happy path" of a `+1` (Proceed) verdict to the more complex scenarios of `-1` (Refuse) and `0` (Sacred Zero). They are essential for understanding the dynamic behavior of the system and for identifying the key decision points and state transitions. The diagrams are based on the interfaces and constants defined earlier and are designed to be a clear and unambiguous representation of the system's logic. They are a valuable tool for developers, auditors, and users who need to understand how the TML framework operates in practice.

The sequence diagrams are divided into three main sections, each corresponding to one of the three possible TML verdicts. The first diagram shows the "Happy Path," where a transaction is evaluated and approved, leading to its execution. The second diagram shows the "Rejection Path," where a transaction is evaluated and found to be in violation, leading to its reversion. The third diagram shows the "Escalation Path," which is the most complex of the three. It illustrates the process of a transaction being put on hold, the summoning of the Stewardship Council, and the eventual resolution of the case. These diagrams are not just a high-level overview; they are intended to be a detailed and accurate representation of the system's behavior, showing the specific function calls, event emissions, and state changes that occur at each step. They are an essential part of the technical documentation for the TML enforcement primitive.

1.3.1. Happy Path: `+1` (Proceed) Transaction

The "Happy Path" sequence diagram illustrates the flow of a transaction that is evaluated and approved by the TML enforcement primitive. This is the most common and straightforward scenario, representing a transaction that is fully compliant with the TML's ethical framework. The diagram begins with a user initiating a transaction by calling a function on the TML-governed contract. The contract then calls the `evaluateTransaction` function, passing the transaction's calldata as an argument. The TML evaluation engine processes the data and returns a `PROCEED` verdict. The contract then emits a `TMLEvaluation` event with the verdict and a unique `logId`, and proceeds to execute the main logic of the function. Finally, the contract emits a `StateTransition` event, although in this case, the state remains `STATE_ACTIVE`. This sequence is designed to be as efficient as possible, minimizing gas costs and latency for legitimate transactions.

The key to the "Happy Path" is the deterministic and efficient nature of the TML evaluation. The evaluation must be thorough enough to ensure that only compliant transactions are approved, but also fast enough to not create a bottleneck for the system. The use of a formal logic system,

such as PLTL, can help to ensure that the evaluation is both rigorous and efficient. The `logId` that is generated and emitted with the `TMLEvaluation` event is also important. It provides a unique identifier that can be used to link the on-chain transaction to off-chain logs and analysis, creating a complete and auditable record of the system's behavior. The "Happy Path" is the ideal scenario, and the TML framework is designed to make this path as smooth and seamless as possible, while still maintaining the highest standards of security and ethical compliance.

1.3.2. Rejection Path: -1 (Refuse) Transaction

The "Rejection Path" sequence diagram shows the flow of a transaction that is evaluated and rejected by the TML enforcement primitive. This path is triggered when a transaction is found to be in clear violation of the TML's ethical axioms. The sequence begins in the same way as the "Happy Path," with a user initiating a transaction and the contract calling the `evaluateTransaction` function. However, in this case, the TML evaluation engine returns a `REFUSE` verdict. The contract then emits a `TMLEvaluation` event with the verdict and a `logId`, and immediately reverts the transaction using the `revert()` statement. This prevents any state changes from occurring and ensures that the system remains secure. The contract remains in the `STATE_ACTIVE` state, and the user is notified that their transaction has been rejected.

The "Rejection Path" is a critical security feature of the TML framework. It provides a robust defense against malicious or non-compliant actions, and it does so in a clear and unambiguous way. The use of the `revert()` statement is important, as it ensures that the transaction is completely rolled back, with no partial state changes. The gas that is consumed up to the point of reversion is still charged, which serves as a deterrent against repeated attempts to violate the rules. The `logId` that is emitted with the `TMLEvaluation` event is also crucial, as it creates a permanent and public record of the violation. This transparency is essential for accountability and for allowing the community to monitor the system's security. The "Rejection Path" is a powerful tool for enforcing the "hard rules" of the TML framework, and it is a key part of what makes the system so secure and reliable.

1.3.3. Escalation Path: 0 (Sacred Zero) Transaction

The "Escalation Path" sequence diagram illustrates the most complex and nuanced flow in the TML framework, triggered by the `0` (Sacred Zero) verdict. This path is taken when the TML evaluation encounters an ethical ambiguity or uncertainty, and it requires human intervention to resolve. The sequence begins with a user initiating a transaction, which is then evaluated by the `evaluateTransaction` function. This time, the function returns a `REVIEW_HOLD` verdict. The contract then emits a `TMLEvaluation` event, transitions to the `STATE REVIEW HOLD` state, and starts a timer. The Stewardship Council is then notified of the pending review, and they begin their deliberations. After a period of discussion and debate, the council reaches a consensus and submits their verdict to the contract. The contract verifies the council's signature and, if the consensus threshold is met, transitions to a new state based on the council's

decision. If the council approves the transaction, the contract executes it and returns to the `STATE_ACTIVE` state. If the council rejects the transaction, the contract reverts it and also returns to the `STATE_ACTIVE` state. If the council fails to reach a consensus within the time limit, the contract may automatically transition to a fail-safe state.

The "Escalation Path" is the cornerstone of the TML governance model. It is the mechanism that allows the system to handle the complex and nuanced ethical questions that cannot be answered by a simple algorithm. The involvement of the Stewardship Council is crucial, as it provides a human element to the decision-making process. The use of a multi-signature scheme and a consensus threshold ensures that the council's decisions are legitimate and representative. The time-bound nature of the review process is also important, as it prevents the system from being stuck in a state of limbo indefinitely. The "Escalation Path" is a sophisticated and well-designed mechanism for bridging the gap between the deterministic world of smart contracts and the nuanced world of human ethics. It is a key part of what makes the TML framework so innovative and powerful.

1.4. Economic Parameters and Gas Cost Estimates

The economic viability of the TML enforcement primitive is a critical consideration for its adoption and long-term sustainability. The cost of interacting with the TML-governed contracts, measured in gas, must be reasonable and predictable for users. The TML framework introduces additional computational overhead compared to a standard smart contract, as every transaction must be evaluated against the TML logic. This overhead must be carefully managed to ensure that the system is not prohibitively expensive to use. The economic parameters of the system, such as the gas costs for different operations and the potential for dynamic pricing based on the contract's state, must be designed to balance the need for security and ethical enforcement with the practical constraints of the EVM's gas market. A thorough analysis of gas costs and the development of optimization strategies are essential for creating a system that is both secure and economically feasible.

The economic model of the TML enforcement layer is based on the principle of "pay for what you use." The gas cost of a transaction is proportional to the amount of computational work that is required to process it. A simple transaction that is immediately approved will have a lower gas cost than a complex transaction that requires a full TML evaluation or even an escalation to the Stewardship Council. This model provides a fair and transparent way to allocate the costs of the system. However, it also means that the gas costs can be volatile, depending on the complexity of the transactions and the congestion of the network. To address this, the TML framework can be designed with a number of optimization strategies, such as using off-chain computation for complex evaluations and implementing a dynamic gas pricing mechanism that adjusts to the current market conditions. The goal is to create a system that is not only secure and ethical but also economically sustainable in the long run.

1.4.1. Gas Cost Analysis for TML Operations

A detailed analysis of the gas costs associated with the various operations of the TML enforcement primitive is essential for understanding its economic impact. The gas cost of a transaction is determined by the sum of the base cost of the transaction and the cost of each opcode that is executed . The TML framework introduces a number of new operations that are not present in a standard smart contract, and each of these operations has a specific gas cost. The following table provides an estimate of the gas costs for the key operations of the TML enforcement primitive. These estimates are based on the current gas schedule of the Ethereum network and are subject to change as the network evolves.

Operation	Description	Estimated Gas Cost	Notes
TML Evaluation	The core logic that evaluates a transaction against the TML axioms.	5,000 - 50,000+	Highly variable depending on the complexity of the axioms and the number of state reads required.
State Transition	Changing the contract's state (e.g., from <code>STATE_ACTIVE</code> to <code>STATE REVIEW_HOLD</code>).	20,000	Primarily the cost of a single <code>SSTORE</code> operation, which is one of the most expensive opcodes .
Event Emission	Logging a <code>TMLEvaluation</code> or <code>StateTransition</code> event.	1,000 - 5,000	The cost depends on the number of indexed topics and the size of the unindexed data.
Merkle Proof Verification	Verifying a Merkle proof against the on-chain root.	3,000 - 10,000	The cost is proportional to the depth of the Merkle tree (i.e., the number of hashes in the proof).
Council Verdict Submission	Submitting and verifying the multi-signature verdict.	10,000 - 20,000	The cost of signature verification using <code>ecrecover</code> for each council member's signature.
Transaction Reversion	Reverting a transaction due to a <code>REFUSE</code> verdict.	2,000 - 10,000	The gas consumed up to the point of reversion, which is typically lower than a successful transaction.

Table 1: Estimated Gas Costs for Key TML Operations

1.4.2. Economic Model for Enforcement Layer

The economic model of the TML system must be designed to be sustainable and to incentivize honest behavior from all participants. For users, the primary cost is the gas fee for their transactions. The model should aim to keep these costs as low as possible while still providing robust ethical enforcement. For the Stewardship Council, there needs to be an incentive to participate in the dispute resolution process. This could be funded through a small fee levied on all transactions, which is then distributed to Council members for their service. Alternatively, the system could be funded by a foundation or a DAO dedicated to promoting ethical AI. The economic model must also consider the cost of the "Always Memory" data architecture. While storing full logs off-chain is cheaper, there are still costs associated with data storage, retrieval, and the generation of Merkle proofs. These costs need to be accounted for in the overall economic design of the system.

1.4.3. Optimization Strategies for Gas Efficiency

To make the TML system economically viable, several gas optimization strategies must be employed. One of the most effective strategies is to minimize on-chain storage. Instead of storing large amounts of data directly in the contract's state variables, the system should use hashes and Merkle roots, as outlined in the "Always Memory" architecture. This significantly reduces the amount of expensive `SSTORE` operations required. Another strategy is to use function modifiers to abstract away common checks, such as state-based access control. This can reduce code duplication and lower the overall bytecode size of the contract, which in turn reduces deployment costs . Using `immutable` and `constant` variables for values that do not change is another effective technique, as these values are embedded directly into the bytecode and do not require a storage read operation to access . Finally, the choice of data types can also impact gas costs. For example, using `uint256` is often more gas-efficient than smaller integer types like `uint8` or `uint16` because the EVM is optimized for 256-bit operations .

1.5. Test Vectors

Comprehensive testing is paramount for the security and reliability of the TML Enforcement Primitive. Given the immutable nature of smart contracts, any bugs or vulnerabilities in the deployed code can have catastrophic consequences, leading to the loss of funds or the compromise of the entire system . Therefore, a rigorous testing suite must be developed to cover all possible execution paths and edge cases. This suite should include unit tests for individual functions, integration tests for the interaction between different components, and end-to-end tests that simulate real-world user scenarios. The test vectors should be designed to validate the correct behavior of the TML logic, the FSM, and the interaction with the Stewardship Council. The use of testing frameworks like Hardhat or Foundry is essential for automating this process and ensuring that the contract behaves as expected before it is deployed to a live network .

1.5.1. Test Cases for +1 (Proceed) Logic

The test cases for the `+1` (Proceed) logic should focus on verifying that the contract correctly allows compliant transactions to execute. This includes testing scenarios where the transaction parameters meet all the ethical requirements defined in the TML axioms. For example, if the TML logic includes a rule that only allows token transfers to whitelisted addresses, a test case would involve calling the transfer function with a whitelisted recipient address. The expected outcome is that the transaction should succeed, the token balance should be updated correctly, and the contract's state should remain `STATE_ACTIVE`. Other test cases should cover the boundaries of the rules. For instance, if there is a limit on the amount of tokens that can be transferred, test cases should be created for transfers just below the limit (which should pass) and at the limit (which should also pass). These tests ensure that the `+1` verdict is correctly applied and that the contract's primary functions work as intended when all conditions are met.

1.5.2. Test Cases for `-1` (Refuse) Logic

The test cases for the `-1` (Refuse) logic are critical for ensuring the system's security and its ability to enforce its ethical rules. These tests should cover all scenarios that are expected to result in a transaction being reverted. Using the same whitelisting example, a test case would involve attempting a transfer to a non-whitelisted address. The expected outcome is that the transaction should be reverted, the token balances should remain unchanged, and the contract should emit a `TransactionReverted` event with a reason code indicating the violation. It is also crucial to test for edge cases and potential attack vectors. For example, test cases should be created to check for integer overflow and underflow vulnerabilities in any arithmetic operations within the TML logic. Reentrancy attacks should also be tested by creating a malicious contract that attempts to re-enter the TML contract during a state update. The goal of these tests is to ensure that the `-1` verdict is a robust and reliable mechanism for preventing non-compliant and malicious actions.

1.5.3. Test Cases for `0` (Sacred Zero) Logic

Testing the `0` (Sacred Zero) logic is the most complex, as it involves the interaction between the TML contract and the Stewardship Council. The test cases should verify that the contract correctly transitions to the `STATE REVIEW HOLD` when the TML logic returns a `0`. This involves creating scenarios that are intentionally ambiguous or not covered by the core axioms. For example, a transaction might involve a new type of interaction that the TML logic is not programmed to handle definitively. The test should verify that the contract emits the `ReviewRequired` event with the correct `logId` and that the contract's state is correctly set to `STATE REVIEW HOLD`. Further tests are needed to simulate the Council's intervention. This can be done by writing test scripts that call the `submitVerdict` function with different outcomes. The tests should verify that the contract correctly processes the Council's verdict and either executes or reverts the original transaction based on the Council's decision. Finally, the time-bounded nature of the `STATE REVIEW HOLD` must be tested. A test case should be created where the review period expires without a verdict from the Council, and the expected outcome is that the original transaction is automatically reverted.

2. The Human Protocol Interface: The Stewardship Council

The Ternary Moral Logic (TML) framework introduces a critical layer of human oversight to counteract the rigid, deterministic nature of smart contract execution. This layer, termed the Human Protocol Interface, is embodied by the Stewardship Council. The Council is not an abstract entity but a concrete software role, implemented as a multi-signature (multi-sig) authority, designed to intervene in specific, well-defined scenarios of ethical ambiguity or uncertainty. Its primary function is to resolve cases where the TML engine outputs a \emptyset (Sacred Zero), indicating a state of review and mandated pause. This section deconstructs the technical architecture of the Stewardship Council, analyzing its role, workflow, and the mechanisms that ensure its power is strictly bounded, preventing it from overriding the core ethical axioms of the TML system. The design ensures that the Council acts as a deliberative body for resolving uncertainty, not as an all-powerful entity that can rewrite the system's fundamental rules.

2.1. Council as a Multi-Sig Authority

The Stewardship Council's power is distributed and safeguarded through a multi-signature (multi-sig) smart contract implementation. This architectural choice is fundamental to preventing unilateral or malicious actions by any single council member. In the Ethereum ecosystem, where native multi-sig support is absent, this functionality is achieved by deploying a dedicated smart contract that governs the approval process for transactions. The TML system leverages this pattern to create a robust and secure governance structure. The multi-sig contract defines a set of n authorized council members (owners) and a threshold m , which specifies the minimum number of signatures required to execute a decision. This "m-of-n" wallet model ensures that no single individual holds absolute power, aligning with the decentralized ethos of blockchain technology and providing a robust defense against both internal bad actors and the loss of private keys. The implementation requires careful auditing, as bugs in the multi-sig contract could lead to frozen or lost funds, a risk mitigated by open-sourcing the code and offering bug bounties.

2.1.1. Role Definition and Responsibilities

The Stewardship Council's role is narrowly defined and procedurally enforced by the smart contract logic. Its sole responsibility is to resolve the `STATE REVIEW HOLD` triggered by a TML evaluation of \emptyset (Sacred Zero). The Council does not have the authority to approve transactions that have been explicitly rejected (-1) by the TML engine. This prohibition is a critical safeguard, ensuring that the Council cannot override the core ethical axioms encoded into the system. Its function is to provide clarity and judgment in situations of ambiguity, not to act as a backdoor for circumventing the rules. The Council's responsibilities can be broken down into several key functions:

- Reviewing Escalated Cases:** When a transaction enters `STATE REVIEW HOLD`, the Council is notified (off-chain) and must review the details of the case, including the TML log ID and the context of the transaction.
- Deliberation and Voting:** Council members deliberate and cast their votes. The multi-sig contract can be designed to support either off-chain signature collection, which is then submitted in a single transaction, or an on-chain confirmation process where members vote individually .
- Executing Verdicts:** Once the `m-of-n` threshold is met, the Council can execute a verdict. This verdict is either to approve the transaction (transitioning the state back to `STATE ACTIVE`) or to reject it (transitioning the state to `STATE FLAGGED` or `STATE INTEGRITY FROZEN` depending on the severity).
- Maintaining Integrity:** The Council is also responsible for maintaining its own integrity. This includes adding or removing members, which itself would require a multi-signature vote, ensuring that the council's composition remains secure and trustworthy .

This structure ensures that the Council's power is not absolute but is instead a carefully managed and distributed responsibility, focused solely on resolving uncertainty within the ethical framework of the TML.

2.1.2. Authentication and Authorization Mechanisms

The authentication and authorization of the Stewardship Council are handled entirely by the underlying multi-signature smart contract. This contract enforces strict access control, ensuring that only authorized council members can participate in the decision-making process. The core of this mechanism is the `onlyOwner` modifier, a common Solidity pattern that restricts function calls to addresses that are registered as owners within the contract . The contract maintains a mapping of owner addresses (`isOwner`) and an array of these addresses (`ownersArr`) to facilitate this check .

There are two primary approaches to implementing the signature collection and verification process within the multi-sig contract :

- Off-Chain Signature Collection:** This approach prioritizes security and simplicity of the on-chain contract. Each council member signs the transaction data (including a `nonce` to prevent replay attacks) off-chain using their private key. These signatures are then collected by an "operator" who submits them all to the contract in a single on-chain transaction. The contract's function then verifies each signature using the `ecrecover()` function, which recovers the signer's address from the signature and compares it against the list of authorized owners. This method minimizes on-chain logic and gas costs for the voting process itself.
- On-Chain Confirmation:** This model pushes more logic onto the blockchain. The contract maintains a list of pending transactions and a mapping of which owners have confirmed each transaction. Any owner can call a `confirmTransaction()` function to cast their vote for a specific proposal. The contract tracks these confirmations, and once

the required threshold is reached, the transaction can be executed. This approach is more transparent, as all confirmations are recorded on-chain, but it requires more on-chain storage and transactions.

For the TML Stewardship Council, a hybrid approach might be optimal. The on-chain confirmation model provides a clear, immutable record of each council member's vote, which is crucial for accountability. However, for efficiency, the final execution could be triggered by a single transaction that includes all the necessary signatures, similar to the off-chain model. The choice of mechanism will depend on the specific trade-offs between on-chain transparency, gas costs, and operational complexity. Furthermore, advanced cryptographic techniques like threshold signatures could be employed to generate a single, compact signature off-chain, which is then verified on-chain, offering a balance between efficiency and security .

2.1.3. Interaction with STATE REVIEW HOLD

The interaction between the Stewardship Council and the `STATE REVIEW HOLD` is the core of the Human Protocol Interface. This interaction is a carefully orchestrated process that ensures the Council's intervention is both necessary and legitimate. When the TML engine evaluates a transaction and outputs `0`, the main TML contract transitions to `STATE REVIEW HOLD`. This state acts as a gate, preventing the transaction from proceeding or being definitively rejected until the Council has had a chance to review it.

The workflow is as follows:

1. **Trigger and Logging:** The TML contract emits a specific event, e.g., `ReviewHoldTriggered(logId, transactionData, timestamp)`, which logs the unique ID of the TML evaluation and the relevant transaction details. This event serves as the official summons for the Stewardship Council.
2. **Off-Chain Notification and Review:** An off-chain service monitoring the blockchain for these events notifies the Council members. They then access the full, encrypted logs (stored off-chain) corresponding to the `logId` to understand the context of the TML's uncertainty.
3. **Council Deliberation:** The Council members use the multi-sig contract to cast their votes. As discussed, this could be through on-chain confirmations or by preparing an off-chain signature.
4. **Verdict Execution:** Once the `m-of-n` threshold is reached, a designated member (or an automated service) calls a function on the TML contract, e.g., `resolveReviewHold(logId, verdict, merkleProof)`. This function takes the original `logId`, the Council's verdict (`+1` to proceed or `-1` to refuse), and a Merkle proof linking the verdict to the on-chain "Always Memory" root.
5. **State Transition:** The TML contract verifies the Merkle proof to ensure the verdict is authentic and corresponds to the correct `logId`. If the proof is valid and the required number of signatures are present, the contract transitions out of `STATE REVIEW HOLD`.

A `+1` verdict returns the state to `STATE_ACTIVE`, while a `-1` verdict could trigger `STATE_FLAGGED` or a permanent freeze, depending on the implementation.

This process ensures that the Council's power is not exercised arbitrarily. It is bound by the cryptographic proofs of the "Always Memory" system and the consensus rules of the multi-sig contract, creating a transparent and accountable mechanism for human-in-the-loop ethical oversight.

2.2. Workflow and Verdict Commits

The operational workflow of the Stewardship Council is designed to be a transparent, auditable, and cryptographically secure process. It begins the moment the TML engine encounters a situation it cannot resolve with a simple `+1` (Proceed) or `-1` (Refuse) and instead outputs a `0` (Sacred Zero). This triggers a multi-step process that involves logging, off-chain review, on-chain voting, and the creation of a cryptographically signed attestation of the final decision. This entire workflow is designed to ensure that the Council's intervention is a deliberate and well-documented act of judgment, not an arbitrary override of the system's logic. The "Verdict Commits" are the central artifact of this process, serving as the immutable, on-chain record of the Council's resolution.

2.2.1. TML Evaluation and Logging

The entire process is initiated by the TML smart contract's evaluation of a transaction. For every transaction that interacts with the contract, the TML engine runs its logic. This logic could involve checking the transaction against a set of ethical rules, consulting an oracle for external data, or evaluating the state of the contract against a set of moral axioms. The result of this evaluation is one of the three TML states: `+1`, `-1`, or `0`.

Crucially, the TML contract does not store the full logs of its decision-making process on-chain, as this would be prohibitively expensive in terms of gas costs and storage. Instead, it implements the "Always Memory" architecture. For each evaluation, the contract performs the following logging actions:

- 1. Emit a Timestamped Event:** The contract emits a detailed event, such as `TMLEvaluated(logId, transactionHash, tmlOutput, timestamp, decisionTreeHash)`. This event provides a high-level summary of the evaluation, including a unique `logId`, the hash of the transaction being evaluated, the TML output (`+1`, `-1`, or `0`), a timestamp, and a hash of the full decision tree or rule set used for the evaluation.
- 2. Update the Merkle Tree (Off-Chain):** The full, detailed logs of the TML evaluation, including the specific rules that were checked and the data that was considered, are stored off-chain in an encrypted format. A new leaf is added to an off-chain Merkle tree, which contains the hash of these detailed logs.

3. **Update the On-Chain Merkle Root:** The root of the off-chain Merkle tree is periodically updated on-chain. This root serves as a cryptographic commitment to the entire history of TML evaluations.

This logging mechanism ensures that while the detailed evidence is kept off-chain for efficiency, a cryptographic proof of its existence and integrity is maintained on-chain. This is essential for the Stewardship Council's review process, as it allows them to verify the context of the TML's `0` output without relying on a trusted third party for the log data.

2.2.2. Council Summoning and Dispute Resolution

When the TML engine outputs a `0`, the contract transitions to `STATE REVIEW HOLD` and emits a specific event, `ReviewHoldTriggered`. This event acts as the "summons" for the Stewardship Council. An off-chain monitoring service detects this event and notifies the council members, providing them with the `logId` of the evaluation in question.

The dispute resolution process then unfolds as follows:

1. **Log Retrieval and Verification:** Council members use the `logId` to retrieve the full, encrypted TML logs from the off-chain storage. They can verify the integrity of these logs by checking their hash against the `decisionTreeHash` emitted in the `TMLEvaluated` event and by verifying the Merkle proof against the on-chain Merkle root.
2. **Deliberation:** The council members review the evidence and deliberate on the correct course of action. This is a human-driven process that occurs off-chain.
3. **Voting:** The council members cast their votes through the multi-sig contract. As previously discussed, this can be done via on-chain confirmations or off-chain signatures.
4. **Consensus and Verdict Formation:** Once the `m-of-n` threshold is reached, a consensus is formed. The final verdict is then prepared for on-chain submission. This verdict is not just a simple `+1` or `-1`; it is a structured piece of data that includes the original `logId`, the final decision, the timestamp of the decision, and the signatures of the council members who voted.

This structured process ensures that the resolution of disputes is not a black box. Every step, from the initial TML evaluation to the final council verdict, is logged and can be audited, providing a high degree of transparency and accountability.

2.2.3. Prohibition on Overriding -1 (Refuse) Verdicts

A cornerstone of the TML architecture is the strict prohibition on the Stewardship Council overriding a `-1` (Refuse) verdict from the TML engine. This rule is not a matter of policy but is enforced at the smart contract level. The TML contract's state machine is designed such that the `STATE REVIEW HOLD` can only be triggered by a `0` output. A `-1` output leads directly to a transaction revert, and the contract does not enter a state where council intervention is possible.

This design principle is crucial for maintaining the integrity of the TML's core ethical axioms. It ensures that the "hard rules" of the system—those that lead to an outright refusal—are immutable and cannot be circumvented by a human authority. The Council's role is to resolve ambiguity (0), not to grant exceptions to the rules (-1).

The rationale behind this prohibition is twofold:

1. **Upholding the "Constitution":** The TML axioms that lead to a -1 verdict are akin to constitutional principles. They represent the non-negotiable ethical foundation of the system. Allowing the Council to override these would be equivalent to allowing a judiciary to unilaterally amend a constitution, undermining the rule of law (or in this case, the rule of logic).
2. **Preventing "God Mode":** This restriction is a key part of the "No God Mode" design. It prevents the Council from becoming a de facto super-administrator with the power to approve any transaction, regardless of the TML's evaluation. This ensures that the system remains trust-minimized and that its security does not rely on the benevolence or incorruptibility of the council members.

By strictly limiting the Council's power to resolving 0 outputs, the TML architecture creates a clear separation of concerns: the deterministic TML engine enforces the rules, and the human Council resolves the edge cases and ambiguities that the engine is not equipped to handle.

2.3. Attestation and Verdict Commits

The final output of the Stewardship Council's deliberation is a cryptographically signed attestation known as a "Verdict Commit." This is not merely a transaction but a formal, on-chain record of the Council's decision, designed to be verifiable, non-repudiable, and inextricably linked to the original TML evaluation that triggered the review. The Verdict Commit serves as the authoritative resolution for the `STATE REVIEW HOLD`, providing the necessary proof for the TML contract to transition to a new state. The architecture of the Verdict Commit is critical for ensuring the integrity and auditability of the entire human-in-the-loop process.

2.3.1. Signing and Verifying Verdict Commits

The creation and verification of a Verdict Commit is a multi-step cryptographic process that leverages the security of the underlying blockchain and the multi-signature contract. The process ensures that the verdict is authentic and has been approved by the required number of council members.

The signing process typically follows the EIP-712 standard, which provides a structured and secure way to sign data off-chain. This standard allows for the creation of a "submessage" that contains the details of the verdict, including:

- The contract address (to prevent replay attacks on other contracts).
- The specific function being called (e.g., `resolveReviewHold`).

- The original `logId` from the TML evaluation.
- The final verdict (+1 or -1).
- A `nonce` to prevent replay attacks.

Each council member signs this structured data with their private key, producing a signature (comprising `v`, `r`, and `s` values). These signatures are then collected.

The verification process occurs on-chain within the TML contract when the `resolveReviewHold` function is called. The function performs the following checks:

1. **Signature Validity:** It uses the `ecrecover()` function to recover the signer's address from each signature and the hashed verdict data. It then checks if the recovered address is a member of the authorized Stewardship Council.
2. **Threshold Check:** It counts the number of valid signatures and ensures that it meets or exceeds the `m-of-n` threshold defined in the multi-sig contract.
3. **Data Integrity:** It verifies that the `logId` and verdict provided in the function call match the data that was signed by the council members.

Only if all these checks pass will the contract accept the Verdict Commit as valid and proceed to update its state. This rigorous verification process ensures that the contract cannot be tricked into accepting a fraudulent or improperly authorized verdict.

2.3.2. Linking Verdicts to TML Log IDs

The link between a Verdict Commit and the original TML evaluation is established through the `logId`. This `logId` is a unique identifier generated by the TML contract for each evaluation and is included in the `ReviewHoldTriggered` event. When the Council prepares its Verdict Commit, the `logId` is a mandatory field in the data that is signed.

This creates an unbreakable cryptographic link:

```
TML Evaluation (logId) -> Review Hold -> Council Deliberation ->
Verdict Commit (includes logId) -> On-Chain Resolution
```

This linkage is crucial for several reasons:

- **Auditability:** It allows anyone to trace a council verdict back to the specific TML evaluation that prompted it. By looking at the `logId` in the Verdict Commit, one can find the corresponding `TMLEvaluated` event and the off-chain logs, providing a complete history of the decision.
- **Preventing Confusion:** It ensures that the Council's verdict is applied to the correct transaction. In a high-throughput system, multiple transactions could be in `STATE REVIEW HOLD` simultaneously. The `logId` prevents a verdict for one transaction from being mistakenly applied to another.

- **Non-Repudiation:** Because the `logId` is part of the signed data, council members cannot later claim that their signature was for a different verdict. The cryptographic link provides undeniable proof of their decision regarding a specific case.

This simple but powerful mechanism of including a unique identifier in both the problem statement (the TML log) and the solution (the Verdict Commit) is fundamental to the integrity and accountability of the human oversight process.

2.3.3. Merkle Proof Verification for State Changes

The final step in the process is the on-chain verification of the Verdict Commit and the subsequent state change. When the `resolveReviewHold` function is called, it must be provided with the Verdict Commit (including the signatures) and a Merkle proof.

The Merkle proof serves to connect the Verdict Commit to the "Always Memory" of the system. The full details of the Council's deliberation, including the final verdict and the signatures, are stored off-chain. A hash of this Verdict Commit is added as a new leaf to the off-chain Merkle tree. The `resolveReviewHold` function requires a Merkle proof that demonstrates that this new leaf (the Verdict Commit) is indeed part of the Merkle tree, the root of which is stored on-chain.

The verification process is as follows:

1. **Verify Signatures:** The contract first verifies the signatures on the Verdict Commit, as described in section 2.3.1.
2. **Verify Merkle Proof:** The contract then takes the hash of the Verdict Commit and the provided Merkle proof and uses a library like OpenZeppelin's `MerkleProof` to verify that this hash is a valid leaf in the Merkle tree committed to by the on-chain root.
3. **Execute State Transition:** If both the signatures and the Merkle proof are valid, the contract is satisfied that the verdict is legitimate and has been properly recorded in the "Always Memory." It then updates its internal state, transitioning out of `STATE REVIEW HOLD`.

This requirement for a Merkle proof adds an extra layer of security. It ensures that the Council's verdict is not just a collection of signatures but is also part of the permanent, tamper-evident record of the system's history. It prevents a scenario where a valid set of signatures could be presented for a verdict that was never officially recorded, further strengthening the "No God Mode" guarantee.

3. Finite State Machine (FSM) Logic & Triggers

The lifecycle of a TML-governed smart contract is formally defined by a Finite State Machine (FSM), a model that ensures predictable and secure behavior by constraining the contract to a finite set of states and governing the transitions between them. This FSM is not merely a

conceptual tool but is directly implemented in the contract's code, making its logic transparent, auditable, and resistant to unintended state changes. The design of this FSM is critical for operationalizing the TML framework, as it provides the structure for enforcing the tri-state logic (+1, 0, -1) and managing interactions with the Stewardship Council. The FSM is composed of four distinct states: `STATE_ACTIVE`, `STATE REVIEW HOLD`, `STATE FLAGGED`, and `STATE INTEGRITY FROZEN`. Each state represents a specific operational mode of the contract, and the transitions between them are triggered by a combination of TML evaluations, Council actions, and time-based conditions. This section provides a detailed mapping of this FSM, including the definition of each state, the precise conditions that trigger state transitions, and the specific, irreversible triggers that lead to the final `STATE_INTEGRITY_FROZEN`, serving as the system's ultimate "kill switch." The FSM approach is a well-established pattern for creating secure smart contracts, with frameworks like FSolidM providing tools to design contracts as FSMs and generate corresponding Solidity code, thereby reducing the potential for semantic gaps and vulnerabilities .

3.1. State Definitions and Transitions

The TML smart contract's operational lifecycle is managed through a carefully designed Finite State Machine (FSM) with four distinct states. This FSM is the core mechanism that enforces the TML's ethical logic and governs the contract's behavior. The use of an FSM is a recognized best practice in smart contract development for creating secure and predictable systems, as it formalizes the contract's behavior and makes it amenable to formal verification . The four states—`STATE_ACTIVE`, `STATE REVIEW HOLD`, `STATE FLAGGED`, and `STATE INTEGRITY FROZEN`—represent the complete set of operational modes the contract can be in. Transitions between these states are not arbitrary; they are triggered by specific, well-defined events, such as the output of the TML evaluation, the consensus of the Stewardship Council, or the detection of a critical security breach. This structured approach ensures that the contract's behavior is always deterministic and aligned with the overarching ethical and security principles of the TML framework. The following subsections will detail the definition and purpose of each of these four states.

3.1.1. STATE_ACTIVE: Standard Operation

`STATE_ACTIVE` is the default and primary operational state of the TML smart contract. When the contract is in this state, it is fully operational and ready to process transactions. All standard functions are available, and the core TML tri-state logic is actively enforced. For any incoming transaction or function call, the contract first invokes the TML evaluation engine. Based on the output of this evaluation (+1, 0, or -1), the contract will then transition to the appropriate next state or execute the requested action. A +1 result leads to the immediate execution of the transaction, keeping the contract in the `STATE_ACTIVE`. A -1 result causes the transaction to revert, also leaving the contract in `STATE_ACTIVE` but with the undesirable action blocked. A 0 result, however, triggers a transition to `STATE REVIEW HOLD`. This state represents the normal, healthy operation of the system, where the automated ethical checks are functioning as

intended. The contract will spend the majority of its lifecycle in this state, processing transactions and making decisions based on the immutable TML logic. The implementation of this state in Solidity would typically involve an `enum` where `ACTIVE` is the initial state, and all primary functions are guarded by a modifier ensuring the contract is in this state before proceeding with the TML evaluation.

3.1.2. STATE_REVIEW_HOLD (Sacred Zero): Escalation State

`STATE_REVIEW_HOLD` is a unique and critical state in the TML FSM, triggered exclusively by the "Sacred Zero" output from the TML evaluation. This state represents a pause in normal operations, initiated when the automated ethical logic encounters a situation of ambiguity or uncertainty that it is not equipped to resolve. It is not a permanent freeze but a temporary holding pattern designed to facilitate human intervention. When the contract enters `STATE_REVIEW_HOLD`, the original transaction is suspended, and the system awaits a resolution from the designated human governance layer—the Stewardship Council. The transition into this state is a deterministic, on-chain event. However, the transition out of this state is contingent on an off-chain action: the Council must reach a consensus and submit a signed verdict to the contract. This design embodies the TML principle that some ethical dilemmas require human judgment. To prevent this state from being exploited as a denial-of-service vector or from causing indefinite delays, it is typically time-bounded. For example, the contract could be programmed to automatically revert the suspended transaction or transition to a `STATE_FLAGGED` status if the Council fails to provide a resolution within a specified timeframe (e.g., 72 hours). This ensures the system remains robust and responsive, even if the human layer is temporarily unavailable or deadlocked. The `STATE_REVIEW_HOLD` is therefore a key feature for integrating human oversight into the automated system in a structured and secure manner.

3.1.3. STATE_FLAGGED: Probationary State

`STATE_FLAGGED` is a probationary state that the contract can enter to signal a heightened level of scrutiny. This state is not triggered by a single TML evaluation but by a pattern of behavior that suggests potential malicious intent or persistent non-compliance. For example, the contract could be programmed to transition to `STATE_FLAGGED` if a specific address or a series of transactions receives a high number of `-1` (Refuse) verdicts within a short period. This state acts as a "yellow card," indicating that the contract is on probation. While in this state, the contract's operations may be subject to additional restrictions. For instance, all transactions might incur a higher gas cost to cover the increased overhead of more intensive auditing. Alternatively, certain functions might be disabled, or all transactions might be automatically escalated to the Stewardship Council for review, even if they would normally receive a `+1` verdict. The `STATE_FLAGGED` state is a flexible tool for managing risk and deterring bad actors without immediately resorting to the nuclear option of `STATE_INTEGRITY_FROZEN`. It provides a graduated response to potential threats, allowing the system to adapt its security posture based on observed behavior.

3.1.4. STATE_INTEGRITY_FROZEN: Irreversible Kill Switch

`STATE_INTEGRITY_FROZEN` is the terminal and irreversible state of the TML smart contract, representing a complete and permanent shutdown of all non-essential functions. This state is the system's "kill switch," designed to be triggered only by the most severe and unambiguous security threats. It is not a state of ethical ambiguity or a tool for routine governance; rather, it is a protective measure of last resort. The triggers for entering `STATE_INTEGRITY_FROZEN` are specific, objective, and verifiable, such as the detection of oracle fraud, evidence of tampering with the on-chain "Always Memory" Merkle root, or unauthorized attempts to access critical administrative functions. This state is irreversible by design, meaning there is no function within the contract's admin surface that can transition the contract back to `STATE_ACTIVE` or any other operational state. This "no-unfreeze" guarantee is a cornerstone of the "No God Mode" principle, providing absolute assurance that even a compromised administrator cannot reactivate a contract that has been frozen for integrity reasons. Once in this state, the contract's primary functions are disabled, and it may only allow for a final, orderly withdrawal of funds or data by their rightful owners, if such a mechanism was pre-programmed. The transition to `STATE_INTEGRITY_FROZEN` is a catastrophic event from an operational perspective, but it is a critical safeguard for protecting the system's core integrity and the assets it holds against fundamental failures or attacks.

3.2. Trigger Conditions for State Transitions

The transitions between the states of the TML FSM are not arbitrary; they are triggered by specific, well-defined conditions that are designed to be objective and verifiable. These triggers can be categorized into three main types: those based on the output of the TML evaluation, those based on the consensus of the Stewardship Council, and those based on time-bound or auto-fail mechanisms. This multi-faceted approach to state transitions ensures that the contract's behavior is both predictable and resilient. It allows the system to respond automatically to the ethical evaluation of transactions, to integrate human judgment in a structured way, and to protect itself from indefinite pauses or deadlocks. The design of these trigger conditions is a critical part of the FSM's logic, as it defines the rules that govern the entire lifecycle of the contract.

3.2.1. TML Output-Based Triggers

The most common triggers for state transitions are the outputs of the TML evaluation logic. As described in the tri-state enforcement mapping, the TML verdict directly dictates the immediate action and the subsequent state of the contract.

- A **+1 (Proceed)** verdict, when the contract is in `STATE_ACTIVE`, results in the execution of the transaction, and the contract remains in `STATE_ACTIVE`. This is the "happy path" and the most frequent transition.

- A **-1 (Refuse)** verdict, when the contract is in `STATE_ACTIVE`, causes the transaction to be reverted, and the contract also remains in `STATE_ACTIVE`. This is a defensive transition that blocks non-compliant actions.
- A **0 (Sacred Zero)** verdict is the most significant TML-based trigger. It causes an immediate transition from `STATE_ACTIVE` to `STATE REVIEW HOLD`, pausing the transaction and escalating it for human review.

These TML-based triggers are the primary drivers of the contract's state machine, ensuring that the ethical evaluation of every transaction has a direct and deterministic impact on the contract's operational state.

3.2.2. Council Consensus-Based Triggers

When the contract is in the `STATE REVIEW HOLD`, the only way to exit this state is through a trigger based on the consensus of the Stewardship Council. This is a critical part of the human-in-the-loop governance model. The trigger condition is the successful verification of a "Verdict Commit" from the Council. This commit must contain a final verdict (**+1** or **-1**) and a sufficient number of valid signatures from authorized council members to meet the predefined quorum threshold (e.g., 5-of-7). Once this consensus is verified on-chain, it triggers a state transition out of `STATE REVIEW HOLD`. If the verdict is **+1**, the contract transitions back to `STATE_ACTIVE` and executes the original transaction. If the verdict is **-1**, the contract may transition to `STATE_ACTIVE` and revert the transaction, or it may transition to `STATE FLAGGED` if the council's decision indicates a serious but not catastrophic issue. This trigger mechanism ensures that the power to resolve ethical uncertainty is not automated but is instead vested in a trusted, human-led governance body.

3.2.3. Time-Bounded and Auto-Fail Mechanisms

To prevent the system from being stuck in a state of limbo indefinitely, the TML FSM incorporates time-bounded and auto-fail mechanisms. The most important of these is associated with the `STATE REVIEW HOLD`. When the contract enters this state, a timer is started. This timer is set to a predefined duration, such as **72 hours**, as specified in the `TMLConstants` library. If the Stewardship Council fails to provide a verdict and reach a consensus within this time period, the timer expires, and this event triggers an automatic state transition. The default action for this auto-fail mechanism is to treat the unresolved ambiguity as a rejection. The contract will automatically revert the suspended transaction and transition back to `STATE_ACTIVE`. This ensures the liveness of the system and prevents a single unresolved dispute from blocking all future operations. This time-bound trigger is a crucial safeguard that balances the need for human deliberation with the practical requirement for a responsive and functional system.

3.3. Integrity Freeze Triggers

The transition to **STATE_INTEGRITY_FROZEN** is the most critical and severe state change within the TML smart contract. It is not triggered by ethical ambiguity or routine governance decisions but is reserved exclusively for objective, verifiable, and catastrophic failures that compromise the fundamental integrity of the system. These triggers are designed to be unambiguous and to represent a clear and present danger to the contract's security or the validity of its data. The "kill switch" is not a tool for punishment or censorship but a protective measure to halt operations when the system's core assumptions are violated. The triggers for this state are carefully defined to prevent its misuse and to ensure it is only activated in genuine emergencies. The following subsections detail the specific conditions that would trigger an integrity freeze, including the detection of oracle fraud, evidence of data tampering, and unauthorized administrative access attempts. These triggers are hard-coded into the contract's logic, making the freeze an automatic and deterministic response to a critical failure.

3.3.1. Oracle Fraud Detection

Many smart contracts rely on external data sources, known as oracles, to provide information that is not available on-chain, such as asset prices or real-world events. The integrity of the TML system can be critically dependent on the trustworthiness of these oracles. If an oracle is compromised and begins to feed fraudulent data to the contract, it can lead to incorrect TML evaluations and potentially catastrophic outcomes. Therefore, a key trigger for **STATE_INTEGRITY_FROZEN** is the detection of oracle fraud. This can be implemented in several ways. The contract could monitor the data from multiple independent oracles and trigger a freeze if there is a significant and persistent discrepancy between them. It could also use a decentralized oracle network like Chainlink, which has built-in mechanisms for detecting and penalizing faulty oracles. If the contract's internal logic detects that the data from a critical oracle is outside of an expected range or is inconsistent with other sources, it should immediately trigger a transition to **STATE_INTEGRITY_FROZEN** to prevent any further actions based on potentially fraudulent information.

3.3.2. Missing or Tampered Log Hashes

A critical trigger for the **STATE_INTEGRITY_FROZEN** state is the detection of a discrepancy between the on-chain "Always Memory" and the off-chain logs it represents. The TML architecture relies on a data model where full logs are stored off-chain for efficiency, while only a cryptographic summary, specifically a Merkle root, is stored on-chain. This on-chain root serves as an immutable and tamper-evident fingerprint of the entire log history. Any sensitive state change or verification process within the contract requires a valid Merkle proof that demonstrates a specific log entry is part of the history summarized by the on-chain root. If, at any point, a required Merkle proof cannot be generated, or if a provided proof fails to match the on-chain root, it is a definitive sign that the off-chain logs have been lost, corrupted, or tampered with. This breaks the fundamental trust assumption of the system. The contract cannot proceed with operations if it cannot verify the integrity of its own history. Therefore, the failure to produce a valid Merkle proof for a required log entry must trigger an immediate and irreversible transition to **STATE_INTEGRITY_FROZEN**. This ensures that the system cannot be manipulated by

altering its historical record and that any action that cannot be cryptographically verified against the on-chain root is blocked, preserving the integrity of the entire system.

3.3.3. Unauthorized Admin Access Attempts

The "No God Mode" principle is a cornerstone of the TML architecture, and a key part of its enforcement is the use of the `STATE_INTEGRITY_FROZEN` kill switch to protect against unauthorized administrative actions. The contract's admin surface is carefully designed and minimized, but it may still contain functions for routine maintenance or upgrades, often protected by `onlyOwner` or similar modifiers. However, the architecture must include a mechanism to detect and respond to unauthorized attempts to access these administrative functions. This could involve monitoring for failed authentication attempts from non-owner addresses, attempts to call admin-only functions without the proper credentials, or any interaction that violates the predefined access control policies. If the contract detects a pattern of such unauthorized access attempts, especially those targeting critical functions that could alter the TML axioms or the state machine logic, it must interpret this as a potential security breach. In such a scenario, the system cannot assume that its administrative layer is secure. To prevent a successful attack from escalating, the contract should be programmed to automatically trigger a transition to `STATE_INTEGRITY_FROZEN`. This proactive defense mechanism ensures that even if an attacker gains access to an admin key, their ability to cause harm is limited, as the system will automatically shut down upon detecting malicious administrative behavior, thus protecting the core integrity of the contract.

4. Data Architecture: The "Always Memory" Implementation

The "Always Memory" data architecture is a cornerstone of the TML smart contract system, designed to reconcile the immutable and transparent nature of the blockchain with the practical and privacy-preserving need to handle complex, potentially sensitive data. This hybrid model strategically partitions data storage between on-chain and off-chain environments. The on-chain component is optimized for verification and integrity checks, storing only cryptographic commitments like Merkle roots and hashes. The off-chain component holds the full, detailed logs, which are encrypted and managed with a proof-of-custody system. This approach ensures that the TML contract can enforce its ethical logic based on a verifiable history of decisions without burdening the network with excessive storage costs or exposing sensitive information. The core principle is that any critical state change within the contract must be substantiated by a valid cryptographic proof, demonstrating that the action is consistent with the logged, off-chain data. This creates a system where trust is not placed in a central authority to maintain the logs, but in the mathematical certainty of cryptographic proofs.

4.1. On-Chain Storage Strategy

The on-chain storage strategy is meticulously designed for efficiency and security, focusing on storing only the essential cryptographic artifacts required for verification. This approach minimizes the contract's storage footprint, thereby reducing gas costs associated with deployment and state updates, a critical consideration for scalable and economically viable smart contracts on EVM-compatible platforms . The primary on-chain data elements are Merkle roots, Keccak-256 hashes of decision trees, and timestamped event emitters. These components work in concert to create a tamper-evident and verifiable record of the TML's operational history. By storing only these compact, fixed-size hashes and roots, the contract can anchor a large volume of off-chain data to the blockchain, ensuring its integrity without incurring the prohibitive costs of storing the raw data itself. This strategy is fundamental to the "Always Memory" concept, providing a lightweight yet robust mechanism for the smart contract to reference and validate its complete operational context.

4.1.1. Storing Merkle Roots

Merkle roots serve as the primary on-chain commitment to the integrity of the off-chain log data. A Merkle tree is a binary tree structure where each leaf node is a hash of a data block (e.g., a specific TML decision log entry), and each non-leaf node is a hash of its two child nodes. The root of this tree, a single 256-bit hash, is stored on-chain. This root acts as a cryptographic fingerprint for the entire set of logs it represents. Any modification to any single log entry would propagate up the tree, resulting in a completely different root hash. This property is crucial for the TML system, as it allows the smart contract to verify the authenticity and immutability of any off-chain log presented to it. When a sensitive state change is requested, the prover must supply a Merkle proof—a path of hashes from the specific log entry up to the stored root. The contract can then recompute the root using this proof; if the computed root matches the stored root, the log is verified as authentic and unaltered. This mechanism is a well-established pattern in blockchain systems for efficiently verifying data integrity without storing the data itself on-chain .

4.1.2. Keccak-256 Hashes of Decision Trees

In addition to Merkle roots, the TML contract stores Keccak-256 hashes of its decision trees. Keccak-256 is the native hash function of the Ethereum Virtual Machine (EVM), implemented as a dedicated opcode, which makes it significantly more gas-efficient than other cryptographic functions like SHA-256 . This efficiency is a key consideration for on-chain operations. The decision trees represent the core logic and axioms of the TML framework. By hashing these trees and storing the resulting `bytes32` hash on-chain, the system creates an immutable reference to the specific ethical ruleset in use. This is critical for the "No God Mode" principle, as it prevents any future upgrades or administrative actions from altering the fundamental TML axioms without detection. Any change to the decision tree would result in a different hash, which would not match the one stored in the contract's immutable storage. This serves as a powerful safeguard, ensuring that the contract's ethical foundation remains constant and verifiable throughout its lifecycle. The use of `keccak256` is a standard practice in Solidity for creating unique identifiers and ensuring data integrity .

4.1.3. Timestamped Event Emitters

Timestamped event emitters are used to create an on-chain, append-only log of significant state transitions and TML evaluations. While the full details of the evaluation are stored off-chain, the contract emits an event every time a TML decision is made (e.g., `TMLDecision(logId, verdict, timestamp)`). These events are stored in the transaction receipts on the blockchain, providing a chronological and immutable record of the contract's activity. The inclusion of a timestamp, derived from the block's timestamp, allows for temporal ordering and auditing. This on-chain log of events is crucial for several reasons. First, it provides a high-level, searchable index of the contract's history, allowing off-chain systems or the Stewardship Council to quickly identify specific decisions or periods of activity. Second, it serves as a public record of the contract's behavior, enhancing transparency and accountability. Third, the `logId` included in the event can be used to link the on-chain event to the corresponding detailed log stored off-chain, creating a verifiable bridge between the two storage layers.

4.2. Off-Chain Storage and Proof-of-Custody

The off-chain storage component of the "Always Memory" architecture is responsible for maintaining the full, detailed logs of all TML evaluations and decisions. This approach is necessary because storing large amounts of data directly on an EVM-compatible blockchain is prohibitively expensive and inefficient. The off-chain system is designed not just for storage but also for security and verifiability. It consists of full encrypted logs, a proof-of-custody mechanism, and a secure access protocol. This ensures that while the data is not on the blockchain, its integrity and availability are guaranteed. The encryption protects the confidentiality of the logs, which may contain sensitive information, while the proof-of-custody mechanism provides a way to prove that the logs have not been tampered with since they were created. This combination of encryption and verifiable custody is essential for building a system that is both private and trustworthy.

4.2.1. Full Encrypted Logs

All detailed logs, containing the full context of TML evaluations (e.g., input parameters, decision tree traversal path, intermediate calculations, and final verdict), are stored off-chain in an encrypted format. This is a critical privacy-preserving measure, as these logs could contain sensitive information about users or the specific ethical dilemmas being evaluated. The encryption ensures that even if the off-chain storage system is compromised, the contents of the logs remain confidential. The encryption keys would be managed by a secure key management system, potentially with access controls that allow only authorized parties, such as the Stewardship Council during a dispute resolution process, to decrypt the logs. The use of strong, industry-standard encryption algorithms (e.g., AES-256) is assumed. This practice of encrypting sensitive off-chain data is a common pattern in decentralized systems that need to balance transparency with privacy, as seen in various off-chain computation and storage solutions .

4.2.2. Proof-of-Custody Mechanisms

To ensure the integrity and availability of the off-chain logs, a proof-of-custody mechanism is implemented. This mechanism provides cryptographic evidence that the logs have been stored correctly and have not been altered. One common approach is to use a system like IPFS (InterPlanetary File System) for decentralized storage. When a log is created, it is encrypted and added to IPFS, which returns a unique content identifier (CID) based on the hash of the content. This CID is then recorded on-chain, either in the event emitter or as part of the data committed to by the Merkle root. Because the CID is a cryptographic hash of the content itself, it serves as a tamper-evident seal. Any change to the log file would result in a different CID, which would no longer match the on-chain record. This provides a strong guarantee of data integrity. Furthermore, the decentralized nature of IPFS ensures that the logs are highly available, as they are replicated across multiple nodes in the network.

4.2.3. Secure Log Access and Retrieval

Access to the off-chain logs must be carefully controlled to maintain privacy and security. A secure access and retrieval protocol is required. This protocol would typically involve a multi-step process. First, a request for a specific log would be made, referencing its on-chain **logId** or CID. Second, the requester's identity and authorization would be verified. For example, only members of the Stewardship Council might be authorized to request logs related to a **STATE REVIEW HOLD**. Third, the encrypted log would be retrieved from the off-chain storage system (e.g., IPFS). Finally, the log would be decrypted using the appropriate key from the secure key management system. This entire process should be logged and auditable to ensure accountability. The protocol must be designed to be resilient against various attacks, such as unauthorized access attempts, man-in-the-middle attacks, and denial-of-service attacks.

4.3. Verification and State Changes

The verification process is the critical link between the on-chain and off-chain components of the "Always Memory" architecture. It is the mechanism by which the smart contract enforces its rules based on the off-chain data. Any sensitive state change, such as resolving a **STATE REVIEW HOLD** or transitioning to **STATE INTEGRITY FROZEN**, must be preceded by a successful verification. This ensures that the contract's state is always a direct and verifiable consequence of the TML's ethical evaluation. The verification process relies on the cryptographic proofs stored on-chain, primarily Merkle proofs, to validate the integrity of the off-chain data presented to it. This creates a system where the contract can be certain that it is acting on authentic and unaltered information, without needing to trust the entity providing the data.

4.3.1. Merkle Proof Verification Process

The Merkle proof verification process is a core function within the TML smart contract. When a user or the Stewardship Council needs to prove the validity of a specific log entry to the contract, they must submit a Merkle proof. This proof consists of the log entry itself, its index in

the tree, and a set of sibling hashes that form a path from the leaf node (the log entry) up to the root. The contract then takes this proof and recomputes the Merkle root. It starts by hashing the log entry. Then, it iteratively hashes this result with each of the provided sibling hashes, moving up the tree. If the final computed hash matches the Merkle root stored in the contract's state, the proof is considered valid. This process is computationally efficient and can be implemented in Solidity with relatively low gas costs, as it primarily involves a series of `keccak256` hash operations.

4.3.2. Requiring Valid Proofs for Sensitive State Changes

The TML contract's state machine is designed to require a valid Merkle proof as a prerequisite for any sensitive state transition. For example, to resolve a `STATE REVIEW HOLD`, the Stewardship Council would need to submit their verdict along with a Merkle proof for the log entry corresponding to the original `0` (Sacred Zero) evaluation. The contract would verify this proof against the stored Merkle root. Only if the proof is valid would the contract accept the council's verdict and transition the state to `STATE ACTIVE` or another appropriate state. Similarly, if a state transition is triggered by an integrity check (e.g., detecting a mismatch in log hashes), the contract might require a proof to confirm the integrity of the "Always Memory" root itself. This hard-coded requirement for cryptographic verification ensures that the contract's state can never be altered based on unverified or potentially fraudulent claims.

4.3.3. Integration with Stewardship Council Verdicts

The verification process is tightly integrated with the workflow of the Stewardship Council. When the council is summoned to resolve a `STATE REVIEW HOLD`, their task is not just to make a decision but to provide a verifiable attestation of that decision. Their "Verdict Commit" must include a reference to the specific TML log ID they are ruling on. To finalize their verdict on-chain, they must submit a transaction that includes their decision and a Merkle proof for the log entry corresponding to that ID. The contract verifies the proof to ensure the council is ruling on a valid, on-record evaluation. This integration ensures that the council's authority is exercised within the constraints of the system's rules. They cannot simply issue a verdict on a non-existent or fabricated event. Their power is not arbitrary; it is grounded in the verifiable history of the TML's own evaluations, creating a system of checks and balances between the automated logic and the human interface.

5. Licensing & Compliance Attestation

The TML smart contract architecture is not just a technical system; it is a framework for enforcing a specific set of ethical principles. To ensure that these principles are aligned with widely accepted standards of human rights and environmental protection, the system includes a mandatory compliance attestation mechanism. This "Genesis" requirement ensures that every TML-governed contract is initialized with a commitment to a specific, verifiable set of ethical documents. This section details the implementation of this compliance layer, including the definition of the "Mandated Corpora," the process for verifying their presence and integrity on

deployment, and the severe consequences for non-compliance. This feature is a key part of the "Logic is Constitution" philosophy, as it embeds the foundational ethical principles directly into the contract's genesis state.

5.1. The "Genesis" Requirement

The "Genesis" requirement is a set of checks and validations that are performed when the TML smart contract is first deployed. This requirement is designed to ensure that the contract is born with a commitment to a specific, immutable set of ethical principles. The core of this requirement is the verification of the "Mandated Corpora," a collection of documents that form the constitutional basis of the TML framework. The contract must not only be aware of these documents but must also be able to verify their integrity, ensuring that they have not been altered or tampered with. This creates a system where the ethical foundation of the contract is not an afterthought but a prerequisite for its very existence.

5.1.1. Mandated Corpora: Human Rights and Earth Protection Protocols

The "Mandated Corpora" is a collection of foundational documents that define the ethical principles of the TML system. This is not a vague or abstract concept but a concrete set of texts that are referenced by the contract. The research prompt specifies that this corpora should include **40+ Human Rights documents** and **26+ Earth Protection protocols**. These documents could include international treaties like the Universal Declaration of Human Rights, the Paris Agreement on climate change, and other widely recognized standards. The specific list of documents would be defined by the TML governance framework and would be represented on-chain by a single, aggregated hash. This approach allows the contract to commit to a large and complex set of principles without needing to store the full text of each document on-chain.

5.1.2. Verifying Presence and Hash of Corpora

On deployment, the TML smart contract must perform a verification of the Mandated Corpora. This is achieved by checking for the presence and integrity of the "Constitution Hash." The "Constitution Hash" is a single, cryptographic hash (e.g., Keccak-256) of the entire Mandated Corpora. This hash would be calculated off-chain by concatenating the hashes of each individual document in a specific order and then hashing the result. The final `bytes32` hash is then passed to the contract's constructor function as a parameter. The constructor function then compares this provided hash against a hard-coded, immutable constant in the contract's code, `TMLConstants.CONSTITUTION_HASH`. If the two hashes match, the contract can be confident that it has been initialized with the correct, unaltered set of ethical principles.

5.1.3. "Constitution Hash" Check on Deployment

The "Constitution Hash" check is the final and most critical step in the "Genesis" requirement. It is performed in the contract's constructor function, which is only executed once during deployment. The logic is simple and absolute:

```
constructor(bytes32 _constitutionHash) {
    require(_constitutionHash == TMLConstants.CONSTITUTION_HASH,
    "TML: Invalid Constitution Hash");
    // If the check passes, the deployment continues.
    // If it fails, the deployment reverts, and the contract is never
    created.
}
```

This check ensures that any attempt to deploy a TML-governed contract with an altered or incorrect set of ethical principles will fail at the very first step. The contract cannot be born without a valid "Constitution Hash," making this a non-negotiable prerequisite for its existence. This is a powerful mechanism for enforcing compliance and ensuring that the entire TML ecosystem is built on a consistent and verifiable ethical foundation.

5.2. Enforcement and Failure Modes

The enforcement of the "Genesis" requirement is absolute and unforgiving. There is no grace period or opportunity for remediation if the Mandated Corpora are found to be missing or altered. The failure mode is immediate and severe, reflecting the critical importance of the ethical foundation of the system. This strict enforcement ensures that the "Logic is Constitution" principle is not just a slogan but a hard-coded reality of the TML architecture.

5.2.1. Defaulting to STATE_INTEGRITY_FROZEN on Non-Compliance

If the "Constitution Hash" check in the constructor fails, the deployment transaction will revert, and the contract will not be created. However, the TML architecture can be designed with an even more severe failure mode. If the contract is deployed successfully but at some later point detects that its foundational axioms have been compromised (e.g., through a malicious upgrade), it can be programmed to transition to `STATE_INTEGRITY_FROZEN`. This would be a last-resort measure to prevent a corrupted contract from continuing to operate. This ensures that the system has a way to protect itself even after deployment, if its core ethical commitments are ever violated.

5.2.2. Immutable Axioms and Rule Sets

The TML axioms and rule sets that are derived from the Mandated Corpora are designed to be immutable. Once the contract is deployed with a valid "Constitution Hash," the specific rules and principles that it enforces are fixed. This immutability is a key feature of the "No God Mode" design, as it prevents any future administrative action from altering the core ethical logic of the system. The contract's behavior is predictable and consistent, as it is always based on the same

set of foundational principles. This provides a high degree of trust and confidence for users and developers who interact with the system.

5.2.3. Preventing Retroactive Alterations

The combination of the "Genesis" requirement and the immutability of the axioms prevents any retroactive alterations of the TML framework. The "Constitution Hash" check ensures that the initial state of the contract is correct, and the immutability of the axioms ensures that it remains correct throughout its lifecycle. This prevents a scenario where a malicious actor could deploy a contract with a valid "Constitution Hash" and then later alter the axioms to serve their own purposes. The TML architecture is designed to be a "constitutional code" where the rules are fixed and cannot be changed, providing a stable and trustworthy foundation for ethical enforcement.

6. Adversarial Analysis & "No God Mode" Proof

A core tenet of the TML architecture is the "No God Mode" principle, which asserts that no entity, including the contract's administrators or the Stewardship Council, should have the power to arbitrarily alter the system's fundamental ethical rules or override its deterministic enforcement. This section provides a detailed adversarial analysis of the system's administrative surface and presents a mathematical proof that the `STATE_INTEGRITY_FROZEN` is an irreversible state, thereby guaranteeing the "No God Mode" property. The analysis differentiates between protective emergency stops and malicious "rug pulls," and it verifies that any upgrade mechanism cannot be used to circumvent the contract's core security guarantees. The goal is to provide a high degree of confidence that the contract's ethical commitments are immutable and that its security mechanisms are irreversible, ensuring that the system operates as a true constitutional code, not a malleable tool for privileged users.

6.1. Admin Surface Analysis

A thorough analysis of the contract's administrative surface is the first step in proving the absence of a "God Mode." This involves a comprehensive enumeration of all functions that have restricted access, typically guarded by modifiers like `onlyOwner` or `onlyAdmin`. The purpose is to map out every possible way a privileged user could interact with the contract's state. This analysis must go beyond a simple code review and consider the underlying patterns and potential vulnerabilities associated with administrative controls in smart contracts. The findings of this analysis will form the basis for the formal proof, as the proof must demonstrate that none of these administrative functions, either individually or in combination, can be used to unfreeze the contract or alter its core axioms.

6.1.1. Enumerating `onlyOwner` and Admin Functions

The first step in the analysis is to systematically identify and document every function in the TML smart contract that is protected by an access control modifier. This includes functions that might

be used for routine maintenance, such as updating oracle addresses or adjusting certain parameters, as well as any functions that could potentially alter the contract's state in a significant way. For each function, its purpose, the conditions under which it can be called, and its potential impact on the contract's state must be clearly documented. This enumeration creates a complete map of the "admin surface," which is the set of all possible actions a privileged user can take. This map is essential for the subsequent security analysis, as it defines the scope of what needs to be proven secure. Any function that can modify the state variables related to the `STATE_INTEGRITY_FROZEN` status or the core TML decision logic must be flagged as a high-priority item for the formal proof.

6.1.2. Mathematical Proof of No Unfreeze Function

The core of the "No God Mode" proof is a mathematical demonstration that no function exists within the TML smart contract that can transition the system from `STATE_INTEGRITY_FROZEN` back to any other state. This is not merely a matter of code review but a formal, mathematical assertion about the contract's state machine. The proof can be constructed by examining the set of all possible state transitions defined in the contract's FSM. Let `S` be the set of all states `{ACTIVE, REVIEW_HOLD, FLAGGED, INTEGRITY_FROZEN}` and `T` be the set of all defined transitions `(s_from, s_to)` where `s_from` and `s_to` are in `S`. The "No God Mode" property can be formally stated as: for all transitions `t` in `T`, if `t.s_from == INTEGRITY_FROZEN`, then no such `t` exists where `t.s_to != INTEGRITY_FROZEN`. In other words, there are no outgoing transitions from the `INTEGRITY_FROZEN` state.

To prove this, one would systematically analyze the code for every function that modifies the contract's state variable (e.g., a function that sets `currentState`). For each such function, we must prove that it is impossible for it to be called when `currentState == INTEGRITY_FROZEN` and have it result in `currentState` being set to a different value. This can be achieved through several mechanisms:

1. **Modifier Guards:** All state-changing functions (except those that transition to `INTEGRITY_FROZEN`) are protected by a modifier like `onlyInState(STATE_ACTIVE)`. This modifier explicitly checks that `currentState` is not `INTEGRITY_FROZEN` and will revert the transaction if it is.
2. **Absence of Logic:** There is simply no code path in the entire contract that allows for a transition from `INTEGRITY_FROZEN` to another state. The functions that set the state are only callable under specific conditions that are mutually exclusive with being in the `INTEGRITY_FROZEN` state.

This proof can be further solidified using formal verification tools and symbolic execution, which can automatically explore all possible execution paths of the contract and verify that no path leads to an "unfreeze" state transition. By providing this formal proof, the TML architecture offers a mathematical guarantee that the kill switch is absolute and that the contract's integrity,

once compromised, cannot be restored by any privileged entity, thus fulfilling the "No God Mode" requirement.

6.1.3. Formal Verification and Symbolic Execution

To provide a rigorous and automated proof of the "No God Mode" property and other security guarantees, the TML smart contract architecture should be subjected to formal verification using symbolic execution and model checking. This approach goes beyond manual code review by using mathematical techniques to formally prove that the contract's behavior conforms to its specification. Frameworks like FSolidM, which generate Solidity code from a formal Finite State Machine (FSM) model, are designed to be integrated with such formal analysis tools . The process begins by creating a formal model of the contract's FSM, which includes its states, transitions, and the conditions that trigger them. This model can then be fed into a symbolic model checker, such as nuXmv, along with a set of properties expressed in a formal specification language like Computational Tree Logic (CTL) .

For the "No God Mode" proof, a key property to verify would be: `AG (currentState == INTEGRITY_FROZEN -> AX currentState == INTEGRITY_FROZEN)`. This CTL formula states that "Globally, it is always the case that if the current state is `INTEGRITY_FROZEN`, then in the next state, the current state will still be `INTEGRITY_FROZEN`." The model checker will exhaustively explore all possible states and transitions of the FSM to verify if this property holds. If it does, the property is proven to be true. If not, the model checker will provide a counter-example—a specific sequence of actions that leads to a violation of the property—allowing developers to identify and fix the flaw in the code. This formal verification process can also be used to check for a wide range of other security vulnerabilities, such as reentrancy, integer overflows, and access control violations, by expressing them as CTL formulae . By incorporating formal verification into its development lifecycle, the TML project can provide a much higher level of assurance in its security and correctness than is possible with traditional testing and auditing alone.

6.2. Timelocks and Emergency Stops

In any robust smart contract system, particularly one with administrative functions, the concepts of timelocks and emergency stops are critical for balancing upgradeability with security. A timelock is a mechanism that introduces a mandatory delay between the announcement of an administrative action (like a contract upgrade or a change in parameters) and its execution. This delay gives users and stakeholders time to review the proposed change and, if they disagree, to withdraw their funds or take other protective measures. An emergency stop, or "circuit breaker," is a mechanism that allows for the immediate suspension of the contract's operations in response to a detected threat. The TML architecture must carefully differentiate between these mechanisms and ensure they are implemented in a way that protects users from malicious actions (like a "rug pull") and ensures that any upgrade mechanism cannot be used to circumvent the contract's core security guarantees.

6.2.1. Differentiating Circuit Breakers from Rug Pulls

The TML contract's `STATE_INTEGRITY_FROZEN` is designed to function as a Circuit Breaker, not a tool for a Rug Pull. The key difference lies in the trigger conditions and the irreversibility of the action. A Circuit Breaker is triggered by objective, verifiable conditions that indicate a security breach or a failure in the system's integrity, such as oracle fraud or a missing log hash. It is a defensive measure designed to protect the system and its users. A Rug Pull, in contrast, is typically triggered by an admin at their discretion and is designed to benefit the admin at the expense of the users. The TML contract's design makes a Rug Pull impossible by ensuring that the `STATE_INTEGRITY_FROZEN` state cannot be triggered by an admin's arbitrary decision and, once triggered, cannot be reversed by anyone. This creates a clear and unambiguous distinction between a legitimate security feature and a potential vector for malicious admin action.

6.2.2. Ensuring Upgrades Cannot Alter TML Axioms

If the TML contract system includes an upgrade mechanism (e.g., a proxy pattern), it is essential to ensure that this mechanism cannot be used to alter the core TML axioms. This is a significant challenge, as upgradeable contracts inherently have a "God Mode" in the sense that the logic can be changed. The TML architecture must mitigate this risk. One approach is to make the core TML decision logic and the `STATE_INTEGRITY_FROZEN` mechanism non-upgradeable. They could be implemented in a separate, immutable contract that the main, upgradeable contract calls. Another approach is to use a timelock contract for all upgrades, requiring a long delay (e.g., several weeks) between when an upgrade is proposed and when it can be executed. This gives users and the community time to review the proposed changes and exit the system if they disagree with them. The upgrade process itself should be subject to a multi-signature governance process, requiring consensus from a diverse group of stakeholders, not just a single admin.

6.2.3. Protecting Against Unauthorized Admin Access

The security of the admin functions is paramount. If an attacker can gain unauthorized access to the admin keys, they could potentially trigger a malicious state change or, in a less secure system, execute a Rug Pull. The TML architecture must implement robust security measures to protect against this. This includes using a multi-signature wallet for the admin account, so that no single key holder can act unilaterally. The threshold for the multi-sig should be set high enough to prevent a small number of keys from being compromised. Additionally, all admin actions should be logged on-chain, creating a transparent and immutable record of all privileged activities. This allows for post-mortem analysis and accountability. Regular security audits of the admin key management process are also essential to ensure that best practices are being followed and that there are no vulnerabilities in the off-chain systems that manage the keys.

7. Conclusion: The Constitutional Code

The technical architecture and governance of TML smart contracts represent a paradigm shift in how we think about decentralized systems. By embedding a deterministic, tri-state ethical evaluation layer directly into the execution logic, the TML framework creates a system where **Ethics Precedes Execution**. This is a fundamental departure from the traditional "Code is Law" paradigm, which often leads to rigid and unforgiving systems that can be exploited or produce unintended negative consequences. The TML architecture, in contrast, is designed to be a "Constitutional Code," where the logic of the system is governed by a higher-order set of ethical principles, not just by the raw instructions of the code. This section summarizes the key architectural principles that enable this shift and explores the broader implications for decentralized governance and the development of ethical AI.

7.1. Ethics Precedes Execution

The core principle of the TML architecture is that every action must be ethically evaluated before it can be executed. This is not an optional feature or a secondary process but the primary gatekeeper for all on-chain activities. The tri-state logic of the TML enforcement primitive ensures that every transaction is met with a clear and decisive response: it is either approved, rejected, or escalated for human review. This creates a system where ethical considerations are not an afterthought but a prerequisite for execution. The "Sacred Zero" state is a particularly important innovation, as it acknowledges the limits of automated reasoning and provides a structured pathway for human judgment to be integrated into the decision-making process. This human-in-the-loop governance model, embodied by the Stewardship Council, ensures that the system can adapt to complex and nuanced ethical dilemmas without compromising its core principles.

7.1.1. Summary of Architectural Principles

The TML architecture is built on a set of core principles that ensure its integrity, security, and ethical alignment. These principles are not just theoretical concepts but are directly implemented in the contract's code, creating a system that is both robust and trustworthy.

- **Deterministic Enforcement:** The TML evaluation logic is a pure function that produces a deterministic output for any given input. This ensures that the contract's behavior is predictable and verifiable by all participants.
- **Tri-State Logic:** The use of a tri-state logic ($+1$, 0 , -1) allows for a more nuanced and sophisticated response to ethical evaluations than a simple binary allow/deny. The "Sacred Zero" state is a key feature that enables human-in-the-loop governance.
- **Finite State Machine:** The contract's lifecycle is governed by a rigorously defined FSM, which ensures that all state transitions are explicit, controlled, and auditable.
- **"No God Mode":** The architecture is designed to prevent any entity, including administrators, from overriding the core ethical axioms or reversing critical security states. The `STATE_INTEGRITY_FROZEN` kill switch is a key component of this principle.

- **"Always Memory":** The hybrid on-chain/off-chain data architecture ensures that the contract's decision-making process is fully auditable and tamper-proof without incurring prohibitive storage costs.
- **Compliance Attestation:** The "Genesis" requirement ensures that every contract is initialized with a commitment to a specific, verifiable set of ethical documents, creating a consistent and trustworthy foundation for the entire ecosystem.

7.1.2. Shift from "Code is Law" to "Logic is Constitution"

The TML architecture represents a fundamental shift from the "Code is Law" paradigm to a new model where **"Logic is Constitution."** In the "Code is Law" model, the rules of the system are defined by the literal instructions of the code, and there is no higher authority to appeal to if the code produces an undesirable or unethical outcome. This can lead to situations where the letter of the law is followed, but the spirit of the law is violated. The "Logic is Constitution" model, in contrast, is based on the idea that the code should be governed by a higher-order set of ethical principles, which are encoded into the contract's logic. The TML axioms, which are derived from foundational documents like human rights charters, serve as the "constitution" of the system. The code is not the law; it is the mechanism for enforcing the law. This shift allows for a more flexible and ethical approach to decentralized governance, where the system can adapt to new challenges and complexities without losing sight of its core values.

7.1.3. Implications for Decentralized Governance and Ethical AI

The TML architecture has significant implications for the future of decentralized governance and the development of ethical AI. By providing a concrete and verifiable mechanism for operationalizing ethical principles, the TML framework offers a new model for building decentralized systems that are not only secure and efficient but also aligned with human values. This could have a profound impact on a wide range of applications, from decentralized finance (DeFi) to supply chain management to social media. The TML architecture also provides a valuable case study for the development of ethical AI. By demonstrating how complex ethical principles can be encoded into a deterministic and verifiable system, the TML framework offers a new approach to the challenge of building AI systems that are both powerful and trustworthy. The "Logic is Constitution" principle could be applied to a wide range of AI systems, from autonomous vehicles to medical diagnosis tools, to ensure that they are governed by a clear and consistent set of ethical principles. The TML architecture is not just a technical solution; it is a new way of thinking about the relationship between technology and ethics, and it offers a promising path forward for building a more just and equitable digital future.